

MTASet: A Tree-based Set for Efficient Range Queries in Update-heavy Workloads

Daniel Manor, Mor Perry, Moshe Sulamy

The Academic College of Tel Aviv-Yaffo



- A concurrent set data structure specifically optimized for environments characterized by high update throughput and frequent range queries
- Based on an (a, b) -tree
- Relaxed balancing
- Lock based
- Optimistic concurrency control
- Leveraging a tailored multi-versioning
- Supports `find(k)`, `insert(k, v)`, `delete(k)` and `scan(fromKey, toKey)`
- Range queries are wait-free
- Linearizable




Figure 2. A snapshot of MTASet, $a=2$, $b=4$. An internal node pointing to a tagged internal node and a leaf node. The tagged internal node points to two leaf nodes. No locks are acquired.




Figure 1. $a=2$, $b=4$. A thread scans a leaf node from left to right, gathering values with the most recent version that is less than or equal to 7. In this scenario, it will collect the values 22 and 903.





Figure 3. $a=2$, $b=4$. An underfull leaf node is being merged with its sibling.

MTASet maintains a **GLOBAL_VERSION** integer variable, which is atomically read and incremented (F&I) by the Scan operation. This version number is used by the scan to determine which values to collect and is read by update operations to assign to the updated values.



Legend: OCC-ABTREE* (blue square), MTASet (red star), Java ConcurrentSkipList (Non-atomic) (green triangle), KiWi (grey diamond), OCC-ABTREE (brown square).